

Noah Anderson

Module 3

6/21/2024

Question 1

```
In [3]: import pandas as pd
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df1
```

```
Out[3]:
```

	Even	Odd
0	2	1
1	4	3
2	6	5

a.

```
In [5]: # Indexing with iloc
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = df1.iloc[[0, 1], [0, 1]]
df2.iloc[0, 0] = 0
print(df2)
print(df1)
```

```

      Even  Odd
0        0    1
1        4    3
      Even  Odd
0        2    1
1        4    3
2        6    5
```

In this case indexing with `iloc` creates a copy of `df1` which explains why `df1` is unchanged when printed.

```
In [7]: # Slicing
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = df1[0:2]
df2.iloc[0, 0] = 0
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	0	1
1	4	3
2	6	5

```
/var/folders/gw/q9_8w8jn0js6k5m5d98vhcdzcx593/T/ipykernel_32314/2128375444.py:4: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df2.iloc[0, 0] = 0
```

In this case slicing creates a view of `df1` which explains why `df1` is changed when printed.

```
In [9]: # Boolean Expression
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = df1[df1.Even < 6]
df2.iloc[0, 0] = 0
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	2	1
1	4	3
2	6	5

```
/var/folders/gw/q9_8w8jn0js6k5m5d98vhcdzcx593/T/ipykernel_32314/1652574861.py:4: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df2.iloc[0, 0] = 0
```

In this case boolean indexing creates a copy when modified as seen by `df1` remaining unchanged when printed.

```
In [11]: # Pandas Methods
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = (df1
      .query('Even < 6')
      .assign(Even = [0, 4])
      )
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	2	1
1	4	3
2	6	5

In this case, indexing through query copies on modify as seen by `df1` being unchanged when printed.

For these use cases, `iloc`, boolean indexing, and querying all copy when indexed whereas slicing modifies in place altering the original data frame.

b.

```
In [14]: # Indexing with iloc using .copy
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = df1.iloc[[0, 1], [0, 1]].copy()
df2.iloc[0, 0] = 0
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	2	1
1	4	3
2	6	5

```
In [15]: # Slicing with .copy
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = df1[0:2].copy()
df2.iloc[0, 0] = 0
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	2	1
1	4	3
2	6	5

```
In [16]: # Boolean Expression .copy
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = df1[df1.Even < 6].copy()
df2.iloc[0, 0] = 0
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	2	1
1	4	3
2	6	5

```
In [17]: # Pandas Methods .copy
df1 = pd.DataFrame({'Even': [2, 4, 6], 'Odd': [1, 3, 5]})
df2 = (df1
      .query('Even < 6')
      .assign(Even = [0, 4])
      ).copy()
print(df2)
print(df1)
```

	Even	Odd
0	0	1
1	4	3

	Even	Odd
0	2	1
1	4	3
2	6	5

All the code chunks have consistent results when using `.copy` now copying on modification. All warnings have disappeared now.

Question 2

a.

```
In [21]: churn = pd.read_csv("churn_modeling.csv")
```

b.

```
In [23]: churn = churn.drop(columns = [ "RowNumber", "Surname", "Gender", "Age"])
```

c.

```
In [25]: churn.isna().any()
```

```
Out[25]: CustomerId      False
        CreditScore    False
        Geography      False
        Tenure         False
        Balance        False
        NumOfProducts  False
        HasCrCard      False
        IsActiveMember False
        EstimatedSalary False
        Exited         False
        dtype: bool
```

With all column values returning `FALSE` for `is.na`, we can conclude that there is no missing data.

d.

```
In [28]: (churn
         .groupby(['Geography', 'Exited']) # Group by Geography and Exited
         .agg(average_estimate = ('Balance', 'mean')) # Calculate mean Balance
         .round(2) # Round to the nearest cent
         )
```

```
Out[28]:
```

		average_estimate
Geography	Exited	
France	0	60339.28
	1	71192.80
Germany	0	119427.11
	1	120361.08
Spain	0	59678.07
	1	72513.35

It does appear that there is a difference between customers who exited and those who did not and this does appear to also vary by country. In Spain and France, those who exited had on average more than \$10,000 those who did not. In Germany the two categories are approximately equal differing only by around \$1,000.

e.

```
In [58]: # Read in new customer data

new_cust = pd.read_csv("new_customers.csv")

# Merge new customer data with the churn data and filter out data where Exit
cust_merge = (pd.merge(new_cust, churn, 'left', 'CustomerId'))
```

```
.query('Exited != 0'))  
  
exited_perc = cust_merge['Exited'].sum() / cust_merge.shape[0] * 100  
  
exited_perc_round = round(exited_perc, 2)  
print("The percent of customers who left and returned to the bank is",  
      exited_perc_round, "%")
```

The percent of customers who left and returned to the bank is 0.86 %

In []: